

Lecture 1d

block 2

Splines


Generalized additive models

Moving beyond simple models

- Sometimes using simple models (linear regression) is not enough
 - Too simple → *more flexible models are needed*

Example: Ambient Assisted Living


- digitally connected and controlled devices for support of people with special needs
- emergency buttons, pressure emergency services with connection to a broader smart home



<https://bstassen.files.wordpress.com/2015/01/aal.png>

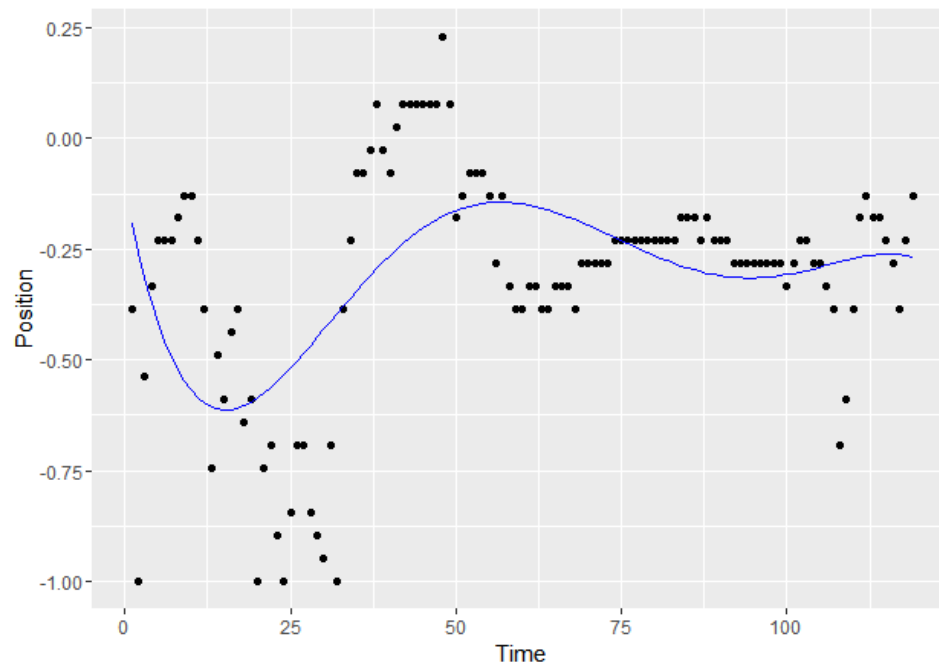
Moving beyond simple models

- Ambient Assisted Living
 - Person's movement is detected by Radio Signal Strength (RSS) measurements → how to remove noise?



Moving beyond simple models

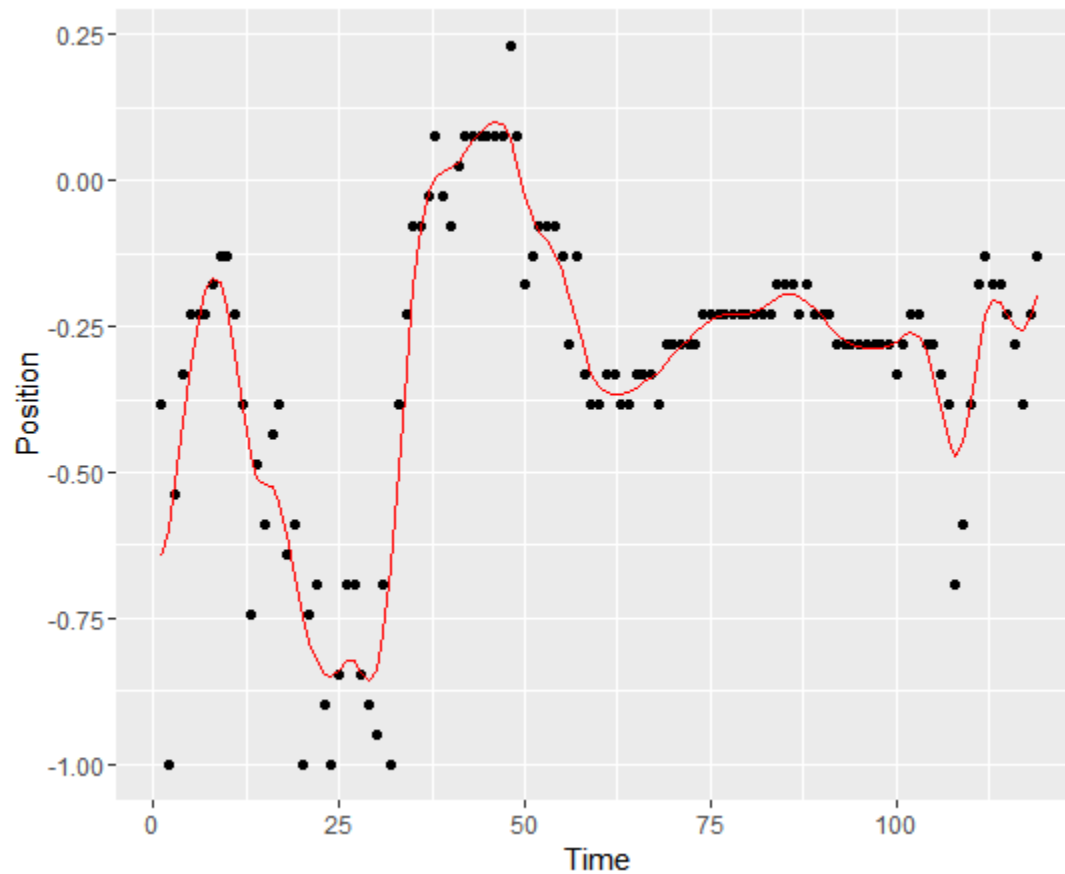
- Attempt 1: 5th degree polynomial



Model underfits data → need more flexible models

Moving beyond simple models

- Using smoothing splines



Basis function expansion

If $y = w_0 + w_1x_1 + w_2x_1^2 + w_3e^{-x_2} + \epsilon$,

Model becomes linear if to recompute:

$$\begin{aligned}\phi_1(x_1) &= x_1 \\ \phi_2(x_1) &= x_1^2 \\ \phi_3(x_1) &= e^{-x_2}\end{aligned}$$


- Any model of the type $Ey = \sum_i w_i \phi_i(\mathbf{x})$ can be fit by linear regression!

Constructing a piecewise linear function

Method A. Introduce linear functions on each interval and a set of constraints

(4 free parameters)

$$\begin{cases} y_1 = \alpha_1 x + \beta_1 \\ y_2 = \alpha_2 x + \beta_2 \\ y_3 = \alpha_3 x + \beta_3 \end{cases}$$
$$\begin{cases} y_1(\xi_1) = y_2(\xi_1) \\ y_2(\xi_2) = y_3(\xi_2) \end{cases}$$



Method B. Use a basis expansion (4 free parameters)

$$h_1(X) = 1, h_2(X) = X, h_3(X) = (X - \xi_1)_+, h_4(X) = (X - \xi_2)_+$$

Theorem. The two methods are equivalent.

Splines


- A piecewise polynomial is called an **order- M** (or degree $M-1$) **spline** if it is continuous and has continuous derivatives up to order $M-2$ at the knots.
- **Equivalent:** An order- M spline with K knots:

$$h_j(X) = X^{j-1}, j = 1, \dots, M$$

$$h_{M+l}(X) = (X - \xi_l)_+^{M-1}, l = 1, \dots, K$$

- An order-4 (degree-3) spline is called a **cubic spline**

In cubic splines, knot discontinuity is not visible



Natural cubic spline

- A cubic spline f is called **natural cubic spline** if its 2nd and 3rd derivatives are zero at a and b


Note that f is linear on extreme intervals

Basis functions of natural cubic splines

$$N_1(X) = 1, N_2(X) = X, N_{k+2} = d_k(X) - d_{K-1}(X), \quad k = 1, \dots, K-2$$

$$\text{where } d_k(X) = \frac{(X - \xi_k)_+^3 - (X - \xi_K)_+^3}{\xi_K - \xi_k}$$

Variance of spline estimators – boundary effects



Fitting smooth functions to data

- Minimize

$$RSS(f, \lambda) = \sum_{i=1}^N (y_i - f(x_i))^2 + \lambda \int \{f''(t)\}^2 dt$$

where λ is **smoothing parameter**.

$\lambda=0$: any function interpolating data

$\lambda=+\infty$: least squares line fit

Optimality of smoothing splines

- The function f minimizing RSS for a given λ is a natural cubic spline with knots at all unique values of x_i (NOTE: N knots!)
- Minimizing sum of squares:

$$f(x) = \sum_{j=1}^N N_j(x) \theta_j = \mathbf{N}(x)^T \boldsymbol{\Theta}$$

$$RSS(\boldsymbol{\Theta}, \lambda) = (\mathbf{y} - \mathbf{N}\boldsymbol{\Theta})^T (\mathbf{y} - \mathbf{N}\boldsymbol{\Theta}) + \lambda \boldsymbol{\Theta}^T \boldsymbol{\Omega}_N \boldsymbol{\Theta}$$

$$\{\mathbf{N}\}_{ij} = N_j(x_i) \quad \{\boldsymbol{\Omega}_N\}_{ij} = \int N_i''(t) N_j''(t) dt$$

$$\hat{\boldsymbol{\Theta}} = (\mathbf{N}^T \mathbf{N} + \lambda \boldsymbol{\Omega}_N)^{-1} \mathbf{N}^T \mathbf{y}$$

A smoothing spline is a linear smoother

- Smoothing spline

$$\hat{f} = \mathbf{N}(\mathbf{N}^T \mathbf{N} + \lambda \Omega_N)^{-1} \mathbf{N}^T \mathbf{y} = \mathbf{S}_\lambda \mathbf{y}$$

is a **linear smoother**.

- Compare with other smoothers, such as linear regression.

Degrees of freedom

- It can be shown that

$$\mathbf{S}_\lambda = (\mathbf{I} + \lambda \mathbf{K})^{-1}$$

where \mathbf{K} is **penalty matrix**

- Eigenvalue decomposition of \mathbf{K} :

$$\mathbf{S}_\lambda = \sum_{k=1}^N \rho_k(\lambda) \mathbf{u}_k \mathbf{u}_k^T$$

$$\rho_k(\lambda) = \frac{1}{1 + \lambda d_k}$$

- d_k and \mathbf{u}_k are eigenvalues and eigenvectors

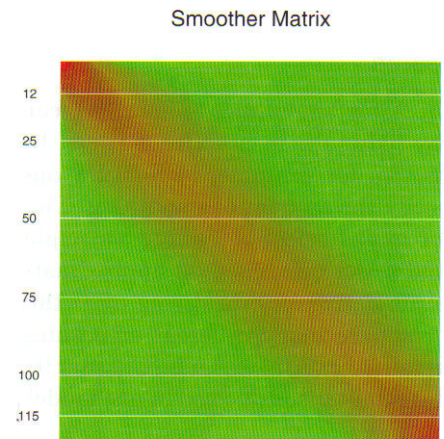
Smoothing splines and shrinkage

$$\mathbf{S}_\lambda \mathbf{y} = \sum_{k=1}^N \mathbf{u}_k \rho_k(\lambda) \mathbf{u}_k^T \mathbf{y}$$

- Smoothing spline decomposes vector \mathbf{y} with respect to basis of eigenvectors and shrinks respective contributions
- The eigenvectors ordered by ρ increase in complexity. The higher the complexity, the more the contribution is shrunk.

Penalty and degrees of freedom

- $df_{\lambda} = \text{trace}(\mathbf{S}_{\lambda}) \rightarrow df_{\lambda} = \sum_{k=1}^N \frac{1}{1 + \lambda d_k}$
- λ increase $\rightarrow df_{\lambda}$ decrease
- higher $\lambda \rightarrow$ higher penalization.
- Smoother matrix is has banded nature
 \rightarrow local fitting method





Automated selection of smoothing parameters

What can be selected:

Regression splines

- Degree of spline
- Placement of knots

Smoothing spline

- Penalization parameter

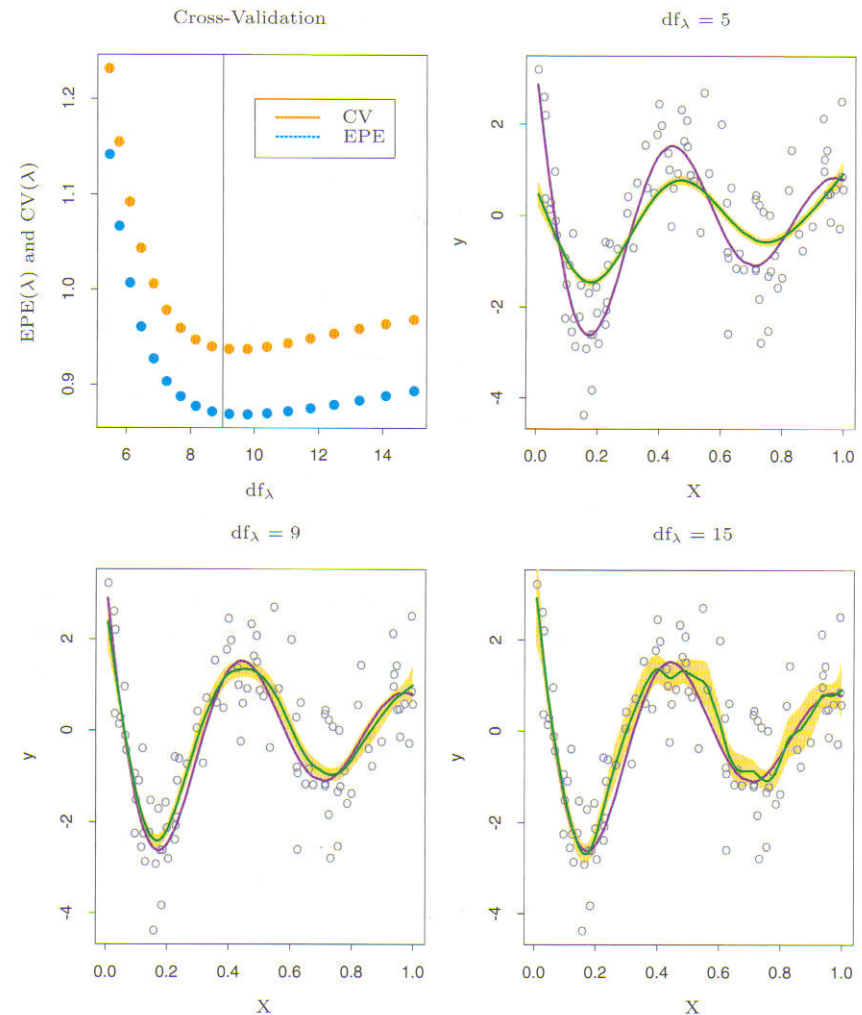
Automated selection of smoothing parameters

$$df_{\lambda} = \text{trace}(\mathbf{S}_{\lambda}) = \sum_{k=1}^N \frac{1}{1 + \lambda d_k}$$

- Use either df_{λ} or λ
 - Given $df_{\lambda} \rightarrow$ solve equation \rightarrow find λ
- Use holdout principle or cross validation for parameter tuning

Automated selection of smoothing parameters

- Bias-variance tradeoff



Multidimensional splines

How to fit data smoothly in higher dimensions?

- Formulate a new problem

$$\min \sum_i (y_i - f(x_i))^2 + \lambda J[f]$$

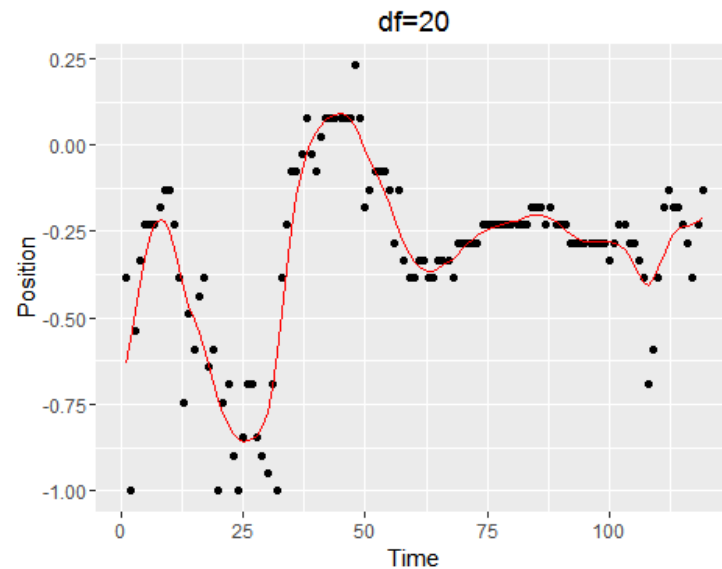

- The solution is **thin-plate splines**
- The solution in 2 dimensions is essentially sum of radial basis functions

$$f(x) = \beta_0 + \beta^T x + \sum \alpha_j \eta(\|x - x_j\|)$$

Splines: R code

- Smoothing splines : `smooth.spline()`
- Natural cubic splines: `ns()` in **splines**
- Thin plate splines: `Tps()` in **fields**

```
res1=smooth.spline(data$Time,data$RSS_anchor2,df=10)  
predict(res1,x=data$Time)$y
```



Generalized additive models

- Model

$$Y \sim EF(\mu, \dots)$$

where

- $g(\mu) = \alpha + s_1(X_1) + s_2(X_2) + s_p(X_p)$
 - $s_i(X)$ - smoothers, normally splines
 - EF – distribution from exponential family
 - g – Link function
- Often linear terms are often included separately

$$EY = \alpha + s_1(X_1) + \dots + s_p(X_p) + \sum_{j=1}^q \beta_j X_{p+j}$$

Example: EF= normal, EF=Bernoulli (logistic)


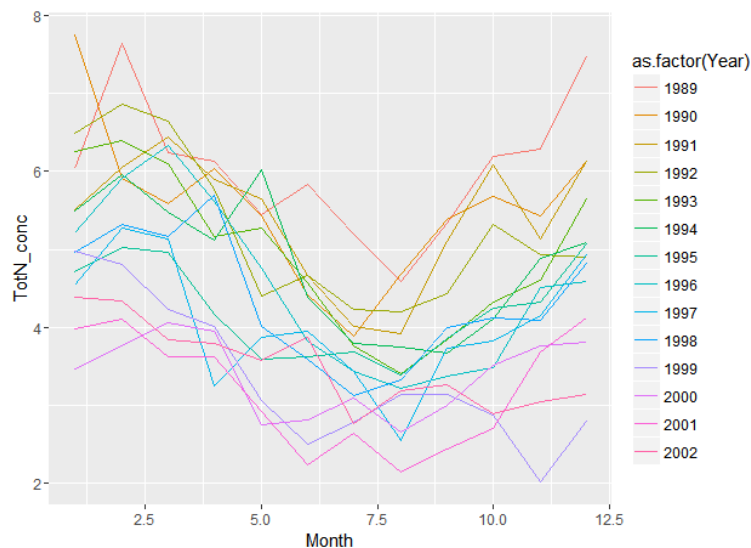
Generalized additive models

- Sometimes even higher orders are included (thin-plate splines)

$$g(\mu) = \alpha + s_1(X_1) + \dots + s_p(X_p) + \sum_{j=1}^q \beta_j X_{p+j} + s_{12}(X_1, X_2)$$

- Method is reasonable to apply when additivity is observed or admissible

Example: Total Nitrogen level in Rhine river



Estimation of additive models

Estimation by MLE

$$g(\mu) = \alpha + f_1(x_1) + \dots + f_p(x_p)$$

The backfitting algorithm for Normal model

1. Initialize: $\hat{\alpha} = \frac{1}{N} \sum_{i=1}^N y_i, \quad \hat{f}_j \equiv 0, \quad j = 1, \dots, p$

2. Cycle: $j = 1, \dots, p, 1, \dots, p, \dots, 1, \dots, p$

$$\hat{f}_j \leftarrow s_j \left[\left\{ y_i - \hat{\alpha} - \sum_{k \neq j} \hat{f}_k(x_{ik}) \right\} \right]$$

$$\hat{f}_j \leftarrow \hat{f}_j - \frac{1}{N} \sum_{i=1}^N \hat{f}_j(x_{ij})$$

λ in each term
can be
estimated by
CV

Generalized additive models

- **Example:** Modelling the concentration of total nitrogen at Lobith on the Rhine
 - There are seasonal trends (GAM reasonable)
 - Variables
 - Nitrogen level
 - Year
 - Month
- R: package **mgcv** (also package **gam**)
 - `gam(formula, family, data, select, method)`
 - Select allows for term (variable) selection
 - `predict()`, `plot()`, `summary()`...
 - `s(k, sp)`
 - k should be the same as the amount of **unique values** of this variable in **smoothing splines**
 - sp - smoothing penalty.

Generalized additive models


- R code

```
library(mgcv)
library(akima)
library(plotly)
river=read.csv2("Rhine.csv")
res=gam(TotN_conc~Year+Month+s(Year,
k=length(unique(river$Year)))+
      s(Month, k=length(unique(river$Month))),
data=river)


s=interp(river$Year,river$Month, fitted(res))
print(res)
summary(res)
res$sp
plot_ly(x=~s$x, y=~s$y, z=~s$z, type="surface")
```

Generalized additive models

λ by CV



$\lambda = 1$



```
> res$sp
      s(Year)      s(Month)
0.0001853275 0.0060877093
```

Generalized additive models

```
> summary(res)
```

```
Family: gaussian  
Link function: identity
```

```
Formula:
```

```
TotN_conc ~ Year + Month + s(Year, k = length(unique(river$Year))) +  
  s(Month, k = length(unique(river$Month)))
```

```
Parametric coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.0006039	0.0016063	0.376	0.707
Year	0.0014321	0.0003277	4.370	2.31e-05 ***
Month	0.2471508	0.1004312	2.461	0.015 *

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Approximate significance of smooth terms:
```

	edf	Ref.df	F	p-value
s(Year)	11.799	12.751	41.32	<2e-16 ***
s(Month)	4.551	5.819	36.51	<2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Rank: 25/27
```

```
R-sq.(adj) = 0.831   Deviance explained = 84.9%
```

```
GCV = 0.26856   Scale est. = 0.23935   n = 168
```

```
> |
```


Generalized additive models

- Seeing trend and seasonal pattern

`plot(res)`

